

Executive Summary

Decision Doc: AI Fluency Rollout - Phased Launch vs. Enterprise Deployment

Overview

HR Ops and L&D has approximately 2.5 FTE of combined project capacity available in H2 FY26 across teams, following the completion of the onboarding redesign project in June. We're evaluating two approaches to launching the company-wide AI fluency program, which has been a stated priority since the January all-hands. This doc presents two options and recommends a phased rollout beginning with managers before expanding to our full workforce, rather than attempting a simultaneous company-wide launch.

Important Considerations

- **Tool access is uneven across the organization.** Engineering and parts of Analytics have had MS Copilot for 6+ months. The majority of the company - roughly 7,200 employees in manufacturing, sales, ops, and admin — has no formally provisioned AI tools and limited informal exposure. Any curriculum that assumes baseline familiarity will fail for most of the audience.
- **Manager readiness is the rate-limiting factor.** Post-training adoption data from the FY25 manager effectiveness series showed that teams whose managers actively modeled new behaviors had 3x the sustained practice rate six weeks out. For AI fluency to stick, managers need to go first. They need enough lead time to actually use MS Copilot before coaching their teams.
- **We have three distinct learner profiles, not one.** Preliminary needs assessment (conducted May–June 2026, n=180 across 12 functions) identified three meaningfully different groups: technical ICs who want depth on prompting and workflow integration; non-technical ICs who want practical, role-specific use cases; and managers who need conceptual grounding plus enough fluency to have credible conversations with their teams. A single curriculum serves none of them well.
- **The L&D team does not have capacity to design three curricula simultaneously.** Current bandwidth allows for one curriculum track in active development at a time. Attempting all three in parallel means all three take longer and none gets adequate iteration before rollout.
- Employee survey data from Q1 FY26 (n=2,400) showed 67% of respondents said they felt behind on AI skills relative to peers at other companies. The pressure to move fast is real. So is the risk of a sloppy launch that generates cynicism instead of adoption.

Assessment

Option 1: Phased rollout — managers first, then ICs in two waves.

Spend Q3 FY26 designing and piloting the manager curriculum (targeting ~800 people managers). Run a live cohort in Q4, collect data, iterate. Launch the technical IC track in Q4 FY26 with the non-technical / ops track following in Q1–Q2 FY27. Total rollout window: 12–15 months.

The case for this approach is operational and behavioral. Managers who complete training 6–8 weeks before their teams create visible champions inside every function. The Q4 manager cohort also gives us a real feedback loop before we're trying to train 9,200 additional people. On the capacity side, designing one track at a time means each track actually gets finished properly before we move to the next.

The downside is timing. Non-technical employees — the largest and least-served group — don't get formal training until FY27. Given the survey data, that gap will generate noise. The communication plan needs to get ahead of it.

Option 2: Design all three tracks concurrently, launch company-wide in Q3–Q4 FY26.

Begin all three curriculum tracks simultaneously, stagger launches by 6–8 weeks (technical ICs first, then managers, then ops/admin), and complete the full company rollout before end of FY26.

The appeal is speed. Everyone gets trained within the fiscal year, the initiative lands as a single visible program, and the 67% of employees who feel behind get something faster. The problem is that doing this requires the L&D team to run three parallel design workstreams while also standing up LMS infrastructure and a new manager-facilitation model. Based on the FY24 DEI training launch, which attempted a similar simultaneous multi-track build under comparable staffing, we ended up with materials that required significant rework post-launch and a net promoter score of 23. Trying to do everything at once under real capacity constraints produces mediocre output across all tracks rather than strong output on any of them.

There's also a sequencing problem embedded in Option 2: if managers receive training at roughly the same time as their direct reports, they're not equipped to reinforce anything. The behavioral case for going managers-first disappears.

Phase / Audience	Headcount	Est. Staff Hours	Timeline
Phase 1: People managers (all levels)	~800	320	Q3 FY26
Phase 2: Individual contributors, technical roles	~4,200	840	Q4 FY26
Phase 3: IC non-technical + ops	~5,000	700	Q1–Q2 FY27
Total	~10,000	1,860	12–15 months

Recommendation

Option 1. The phased approach is slower by design, and that's the point. A company-wide AI fluency program that generates cynicism or becomes a checkbox exercise actively sets back the goal. Going managers-first gives us a feedback loop, builds internal champions, and keeps the L&D team building one thing well at a time instead of three things poorly at once. The non-technical ops population gets a more informed, iterated curriculum for having waited two quarters.

The one thing we're asking for to make this work: explicit executive air cover for the 12–15 month timeline. The Q1 FY26 survey data will resurface, employees will ask why ops teams are last, and managers will need to be able to point to something official rather than ad-hoc explanations.

People & Learning is requesting approval to begin the manager curriculum build in August, with a target pilot cohort date of November 2026, and a full company-wide completion target of Q2 FY27.

Reviewed By: Sarah Okonkwo, Marcus Delgado, Priya Anand

APPENDIX

Current AI Tool Access

Employee Segment	Headcount	Current AI Tool Access
Engineering / Product	~1,400	GitHub Copilot, MS Copilot, Cursor
Analytics / Data	~600	MS Copilot
Managers (all functions)	~800	MS Copilot
Sales / Customer Success	~2,200	MS Copilot, Salesforce Einstein
Manufacturing / Operations / Admin	~5,000	None formally provisioned

Option 1 — Phased Rollout: Year-One Capacity Estimate

Workstream	Owner	Est. Staff Hours	Timeline
Curriculum design (3 tracks)	L&D + Comms	210	Q2–Q3 FY26
LMS configuration and vendor mgmt	IT + L&D	160	Q2 FY26
Manager facilitation training	L&D	90	Q3 FY26

Comms, launch, and change mgmt	Comms	120	Q3–Q4 FY26
Compliance tracking and reporting	HR Ops	80	Ongoing
Total		730	6–9 months

Risk Notes

- **Technical IC resistance to "basic" content.** Developers who've been using Copilot for a year will disengage from a 101-level curriculum. The technical track needs a genuine intermediate tier, not just a renamed version of the manager content. This is solvable with the right SME input during curriculum design; it requires building in that time.
- **Manager time commitment.** Asking 800 managers to complete an 8–10 hour program in Q4 alongside year-end cycles is a real ask. Recommend a self-paced design with a single 90-minute live session, not a cohort-based model. The data from the FY25 manager series supports this — completion rates for self-paced with one live anchor were 74% vs. 41% for cohort-only.
- **Vendor selection.** Three vendors were scoped in April (Udemy for Business, AI Learning Co., and an internal build option). Internal build is off the table at current headcount. A vendor decision needs to happen by the end of July to hit the August curriculum build start. IT and Procurement need to be looped in by June 30.
- **The "why are ops teams last" question.** This will come up. The answer is straightforward - managers going first improves outcomes for everyone, and we're building a better curriculum for ops by iterating first - but it needs to be in the official communications, not discovered by employees through word of mouth.